


Knowledge Graphs for Enhanced Cross-Operator Incident Management and Network Design

draft-tailhardat-nmop-incident-management-noria-02



IETF NMOP interim meeting - 21st, May 2025

Lionel TAILHARDAT, Orange Research, lionel.tailhardat@orange.com


Fano RAMPARANY, Orange Research, fano.ramparany@orange.com

Pauline FOLZ, Orange Research, pauline.folz@orange.com

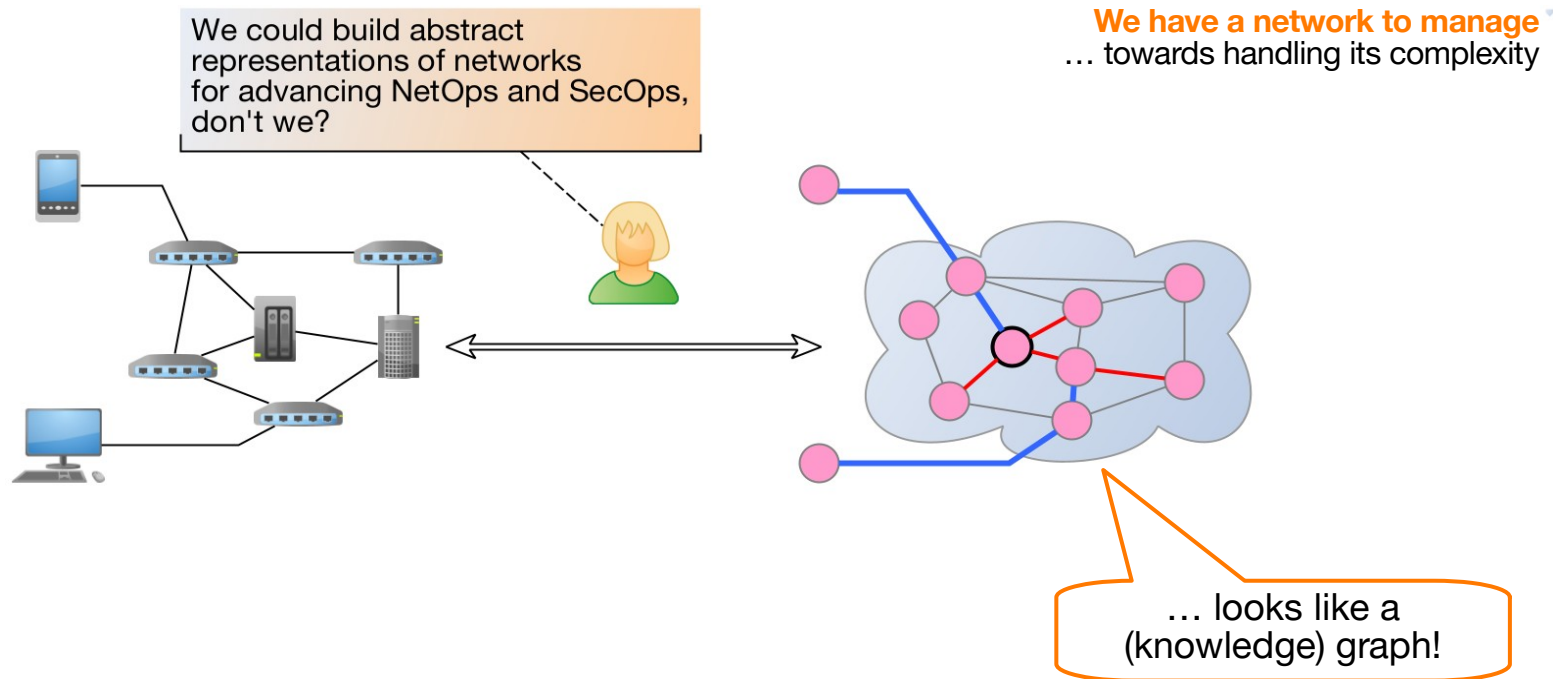


Context & motivations: abstracting networks & sharing behavioral models

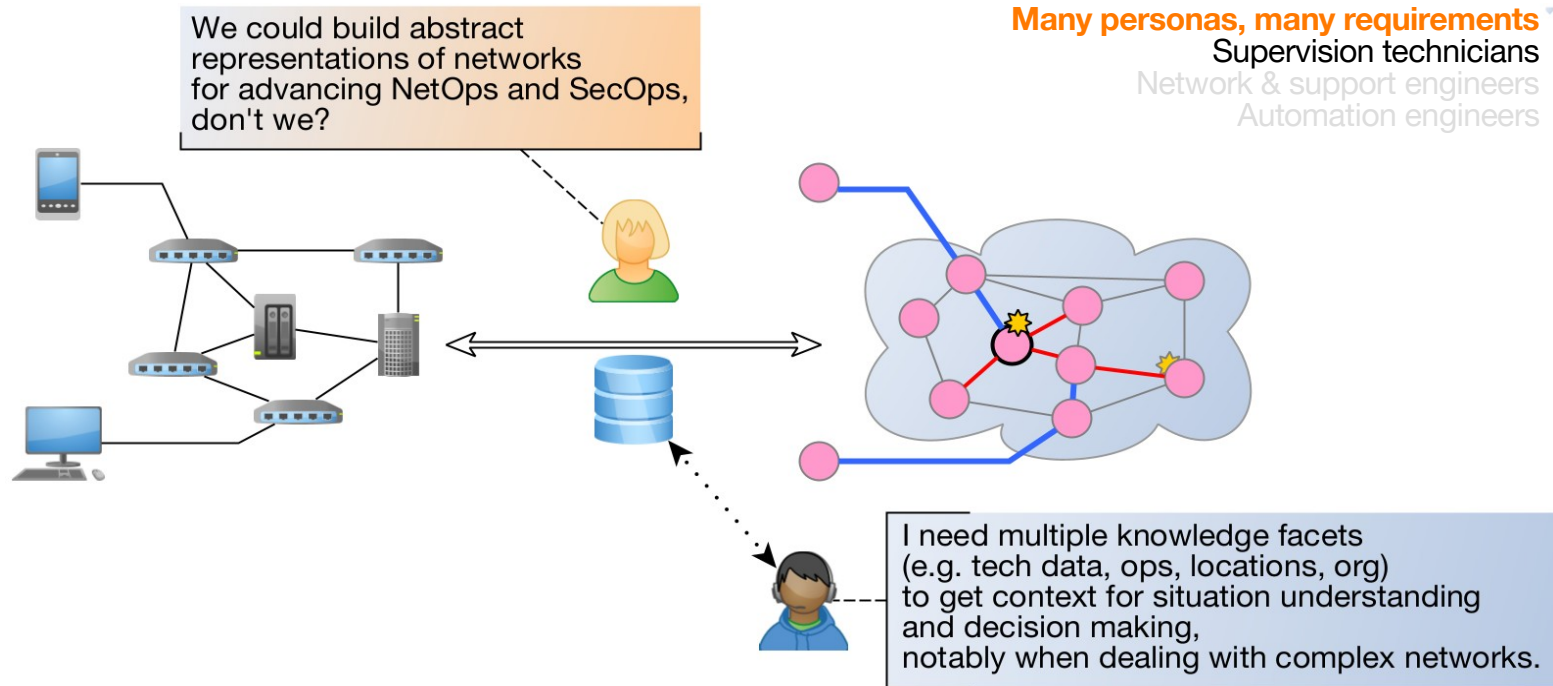
We have a network to manage
... towards handling its complexity



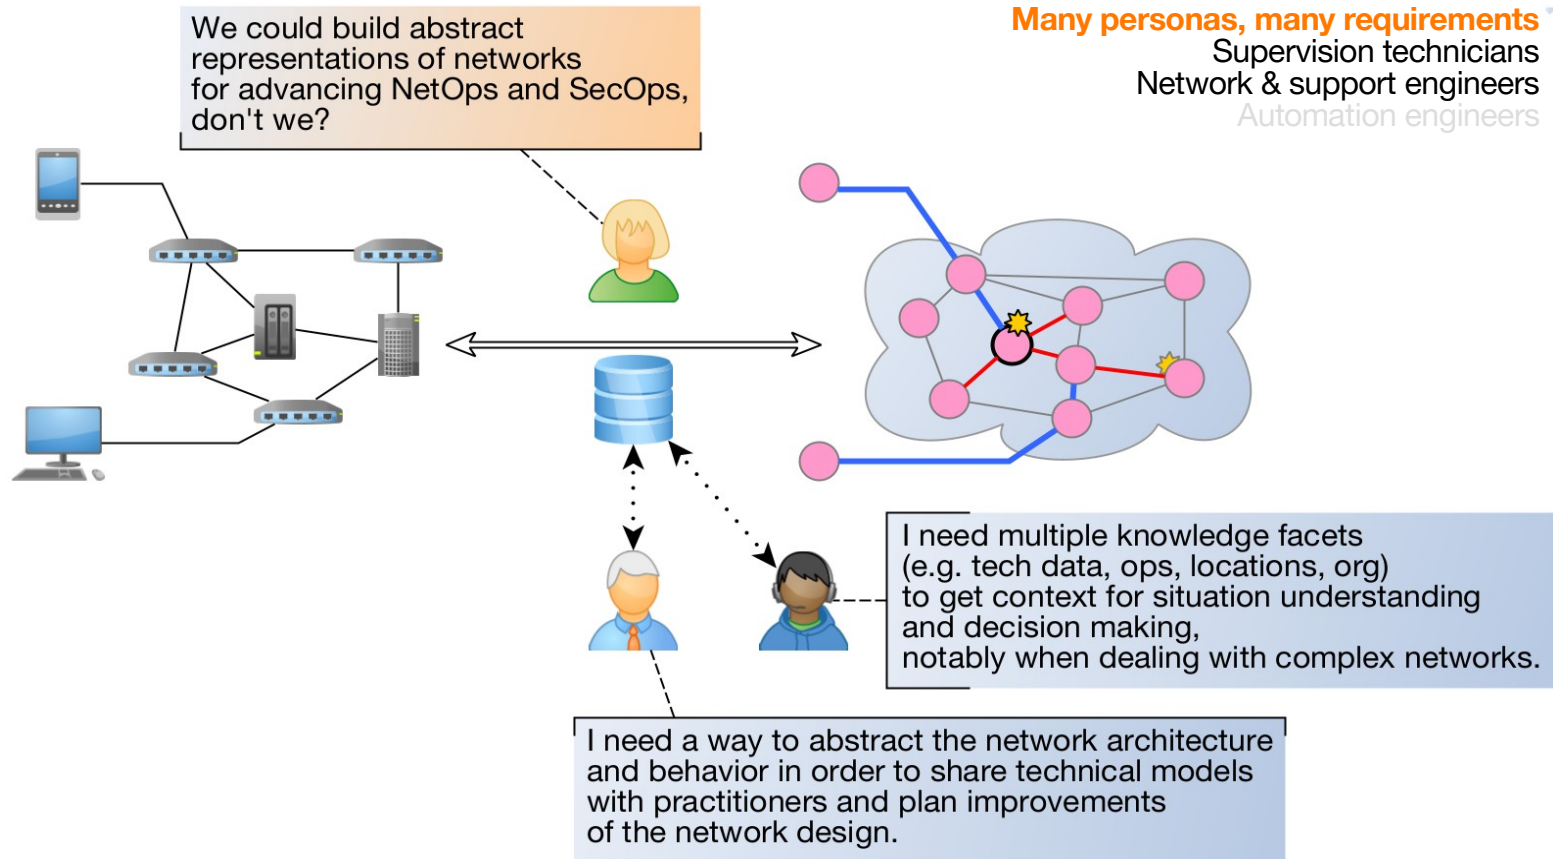
Context & motivations: abstracting networks & sharing behavioral models



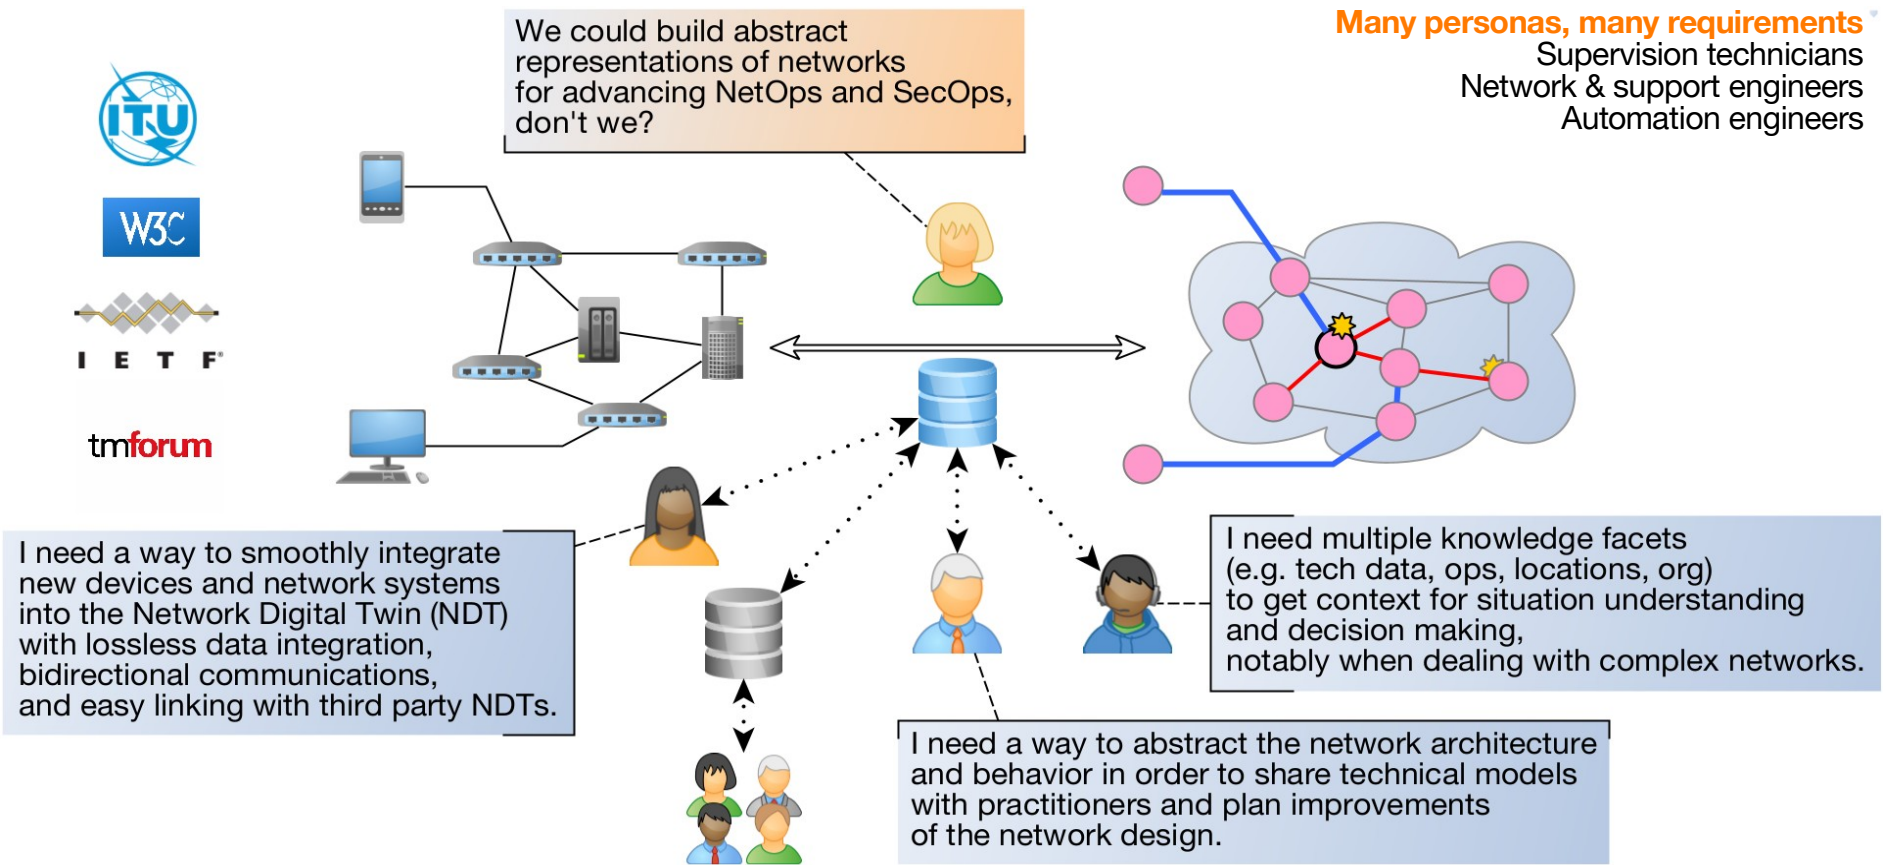
Context & motivations: abstracting networks & sharing behavioral models




Context & motivations: abstracting networks & sharing behavioral models



Context & motivations: abstracting networks & sharing behavioral models



Context & motivations: abstracting networks & sharing behavioral models



IETF NMOP interim meeting n°3, 2024-09-11
draft-tailhardat-nmop-incident-management-noria-01
[I-D, presentation] ... raised the ideas that:

Data Knowledge graph as a combination of a Digital Map [I-D] with operational data and Operational Support Systems (OSS) data.

Opportunity YANG-based configuration data can be converted to build a Digital Map, thereby connecting the Decision Support Systems (DSS) with network production.

draft-tailhardat-nmop-incident-management-noria: overview

Workgroup: Network Management Operations
Internet-Draft: draft-tailhardat-nmop-incident-management-noria-latest
Published: 15 May 2025
Intended Status: Informational
Expires: 16 November 2025
Authors: L. Tailhardat, R. Troncy, Y. Chabot
Orange Research, EURECOM, Orange Research
F. Ramparany, P. Folz
Orange Research, Orange Research

Knowledge Graphs for Enhanced Cross-Operator Incident Management and Network Design

Abstract

Operational efficiency in incident management on telecom and computer networks requires correlating and interpreting large volumes of heterogeneous technical information. Knowledge graphs can provide a unified view of complex systems through shared vocabularies. YANG data models enable describing network configurations and automating their deployment. However, both approaches face challenges in vocabulary alignment and adoption, hindering knowledge capitalization and sharing on network designs and best practices. To address this, the concept of a IT Service Management (ITSM) Knowledge Graph (KG) is introduced to leverage existing network infrastructure descriptions in YANG format and enable abstract reasoning on network behaviors. The key principle to achieve the construction of such ITSM-KG is to transform YANG representations of network infrastructures into an equivalent knowledge graph representation, and then embed it into a more extensive data model for Anomaly Detection (AD) and Risk Management applications. In addition to use case analysis and design pattern analysis, an experiment is proposed to assess the potential of the ITSM-KG in improving network quality and designs.

Table of Contents



- 1. Introduction
- 2. Conventions and Definitions
- 3. An ITSM-KG for Learning and Sharing Network Behavioral Models
 - 3.1. Principles
 - 3.2. Relation to the Digital Map
 - 3.2.1. Core Requirements
 - 3.2.2. Design Requirements
 - 3.2.3. Architectural Requirements
- 4. Strategies for the ITSM-KG Construction
 - 4.1. From YANG-based Configurations to Meta-Knowledge Graph
 - 4.2. Implementing Alignments of Model-Specificities to a Multi-Faceted Knowledge Graph
 - 4.2.1. The Network of Ontologies Approach
 - 4.2.2. Explicit Linking in the ONTO-META
 - 4.3. Extract-Transform-Load Pipelines for the ITSM-KG
 - 4.3.1. Handling Event Streams
 - 4.3.2. Federated Data Architecture
- 5. Experiments
 - 5.1. Experimental Plan
 - 5.2. Implementation Status
 - 5.2.1. NORIA
 - 5.2.2. YANG2OWL
- 6. Security Considerations
- 7. IANA Considerations
- 8. References
 - 8.1. Normative References
 - 8.2. Informative References

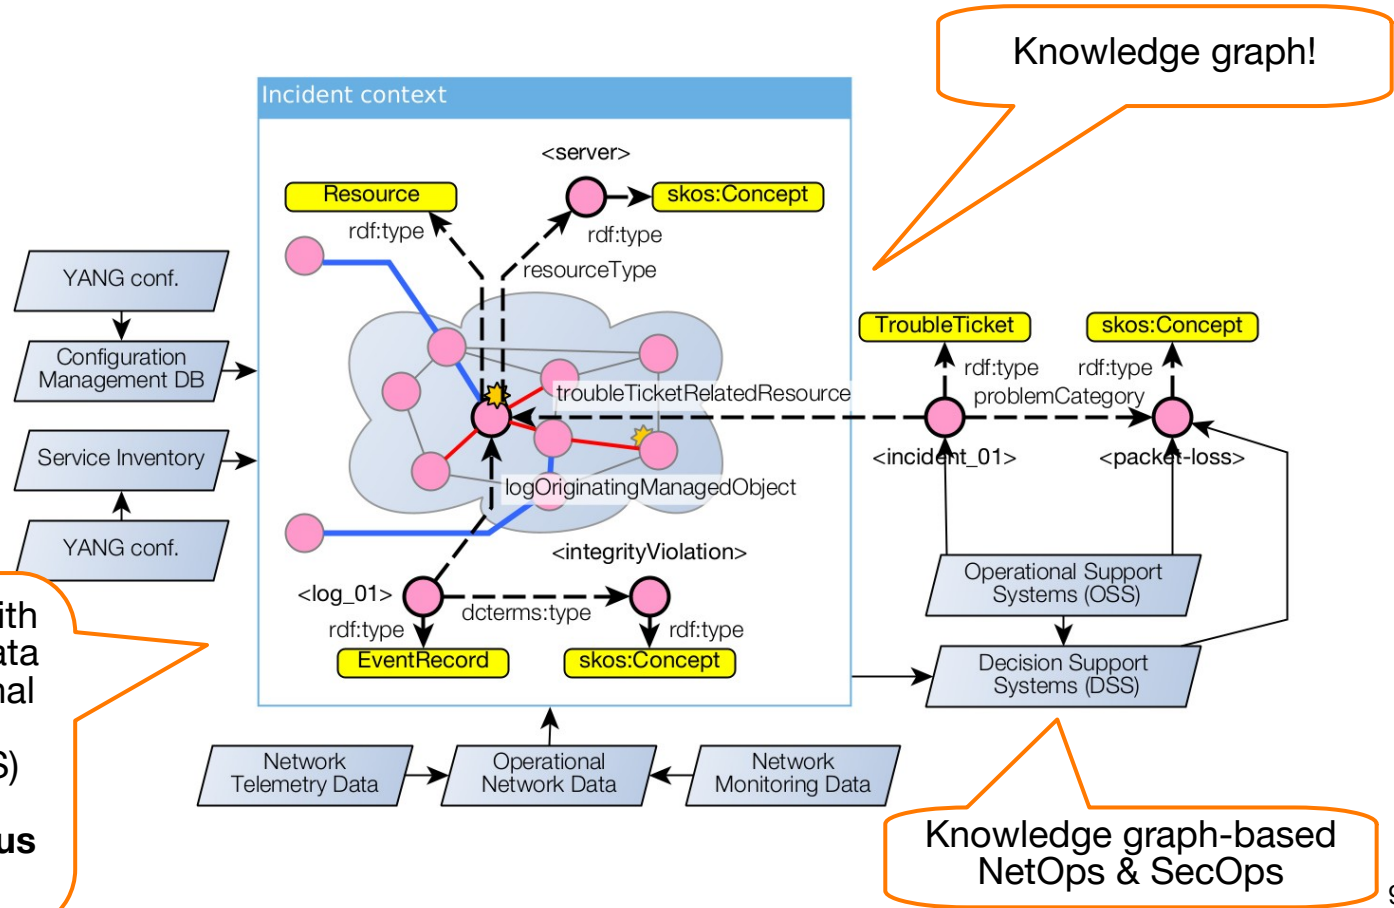
Acknowledgments

draft-tailhardat-nmop-incident-management-noria: overview -- §3

Table of Contents

- 1. Introduction
- 2. Conventions and Definitions
- 3. An ITSM-KG for Learning and Sharing Network Behavioral Models
 - 3.1. Principles
 - 3.2. Relation to the Digital Map
 - 3.2.1. Core Requirements
 - 3.2.2. Design Requirements
 - 3.2.3. Architectural Requirements
- 4. Strategies for the ITSM-KG Construction
 - 4.1. From YANG-based Configurations to Meta-Knowledge Graph
 - 4.2. Implementing Alignments of Model-Specificities to a Multi-Faceted Knowledge Graph
 - 4.2.1. The Network of Ontologies Approach
 - 4.2.2. Explicit Linking in the ONTO-META
 - 4.3. Extract-Transform-Load Pipelines for the ITSM-KG
 - 4.3.1. Handling Event S...
 - 4.3.2. Federated Data
- 5. Experiments
 - 5.1. Experimental Plan
 - 5.2. Implementation Sta
 - 5.2.1. NORIA
 - 5.2.2. YANG2OWL
- 6. Security Considerations
- 7. IANA Considerations
- 8. References
 - 8.1. Normative Reference
 - 8.2. Informative Referen
- Acknowledgements

Digital Map with operational data and Operational Support Systems (OSS) data = heterogeneous data sources



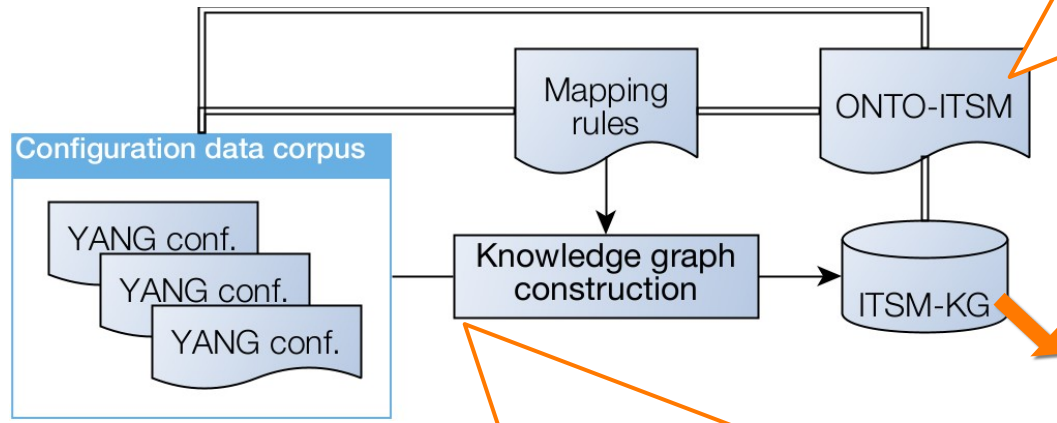
draft-tailhardat-nmop-incident-management-noria: overview -- §4.1

Table of Contents

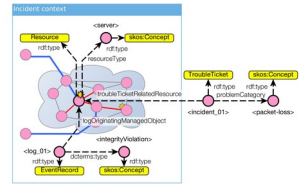
- 1. Introduction
- 2. Conventions and Definitions
- 3. An ITSM-KG for Learning and Sharing Network Behavioral Models
 - 3.1. Principles
 - 3.2. Relation to the Digital Map
 - 3.2.1. Core Requirements
 - 3.2.2. Design Requirements
 - 3.2.3. Architectural Requirements
- 4. Strategies for the ITSM-KG Construction
 - 4.1. From YANG-based Configurations to Meta-Knowledge Graph
 - 4.2. Implementing Alignments of Model-Specificities to a Multi-Faceted Knowledge Graph
 - 4.2.1. The Network of Ontologies Approach
 - 4.2.2. Explicit Linking in the ONTO-META
 - 4.3. Extract-Transform-Load Pipelines for the ITSM-KG
 - 4.3.1. Handling Event Streams
 - 4.3.2. Federated Data Architecture
- 5. Experiments
 - 5.1. Experimental Plan
 - 5.2. Implementation Status
 - 5.2.1. NORIA
 - 5.2.2. YANG2OWL
- 6. Security Considerations
- 7. IANA Considerations
- 8. References
 - 8.1. Normative References
 - 8.2. Informative References

Acknowledgements

ONTO-ITSM Brings a unified view of the network and its ecosystem, i.e. enables querying/traversing the ITSM-KG with a shared vocabulary.



IT Service Management Knowledge Graph



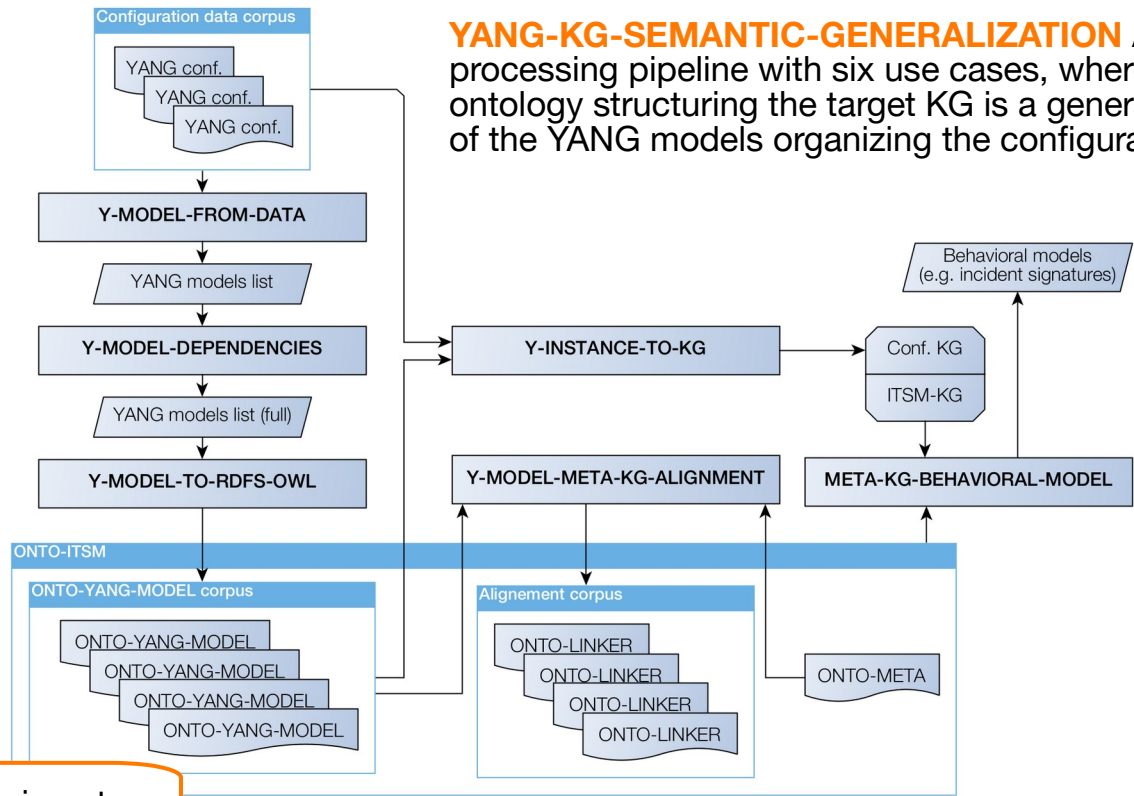
YANG conf. The YANG language is the current standard in network management and will remain so. Including this data—without loss of information and while respecting each operator's specifics (e.g. choice of YANG modules)—can be done by reflecting on **how to project YANG configuration data into the broader ITSM-KG discourse domain.**

draft-tailhardat-nmop-incident-management-noria: overview -- §4.2 & §5.1

Table of Contents

- 1. Introduction
- 2. Conventions and Definitions
- 3. An ITSM-KG for Learning and Sharing Network Behavioral Models
 - 3.1. Principles
 - 3.2. Relation to the Digital Map
 - 3.2.1. Core Requirements
 - 3.2.2. Design Requirements
 - 3.2.3. Architectural Requirements
- 4. Strategies for the ITSM-KG Construction
 - 4.1. From YANG-based Configurations to Meta-Knowledge Graph
 - 4.2. Implementing Alignments of Model-Specificities to a Multi-Faceted Knowledge Graph
 - 4.2.1. The Network of Ontologies Approach
 - 4.2.2. Explicit Linking in the ONTO-META
 - 4.3. Extract-Transform-Load Pipelines for the ITSM-KG
 - 4.3.1. Handling Event Streams
 - 4.3.2. Federated Data Architecture
- 5. Experiments
 - 5.1. Experimental Plan
 - 5.2. Implementation Status
 - 5.2.1. NORIA
 - 5.2.2. YANG2OWL
- 6. Security Considerations
- 7. IANA Considerations
- 8. References
 - 8.1. Normative References
 - 8.2. Informative References

Acknowledgments



YANG-KG-SEMANTIC-GENERALIZATION A data processing pipeline with six use cases, where the ontology structuring the target KG is a generalization of the YANG models organizing the configuration data.

Call for experiments

draft-tailhardat-nmop-incident-management-noria: overview -- §4.3.1

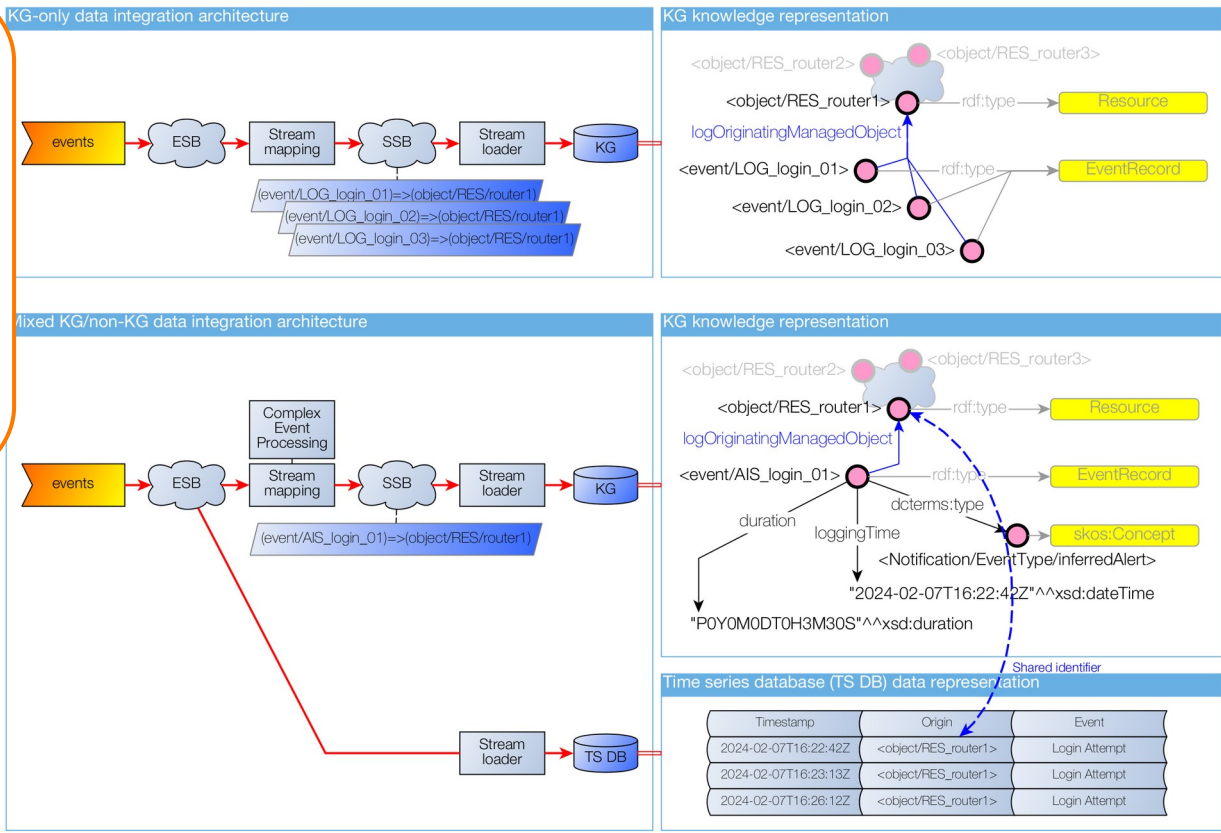
Handling Event Streams Scenarios for constructing a ITSM-KG through an Extract-Transform-Load (ETL) data integration pipeline.

Event streams can be high-paced: it could be beneficial to **leverage input/output (I/O) performance optimizations specific to each type of database management system (DBMS)**, such as Time-Series DataBases (TSDBs) for streaming data and graph databases for knowledge graphs.

- 4.2.6. METX
- 4.3. Extract-Transform-Load Pipelines for the ITSM-KG
 - 4.3.1. Handling Event Streams
 - 4.3.2. Federated Data Architecture

- 5. Experiments
 - 5.1. Experimental Plan
 - 5.2. Implementation Status
 - 5.2.1. NORIA
 - 5.2.2. YANG2OWL
- 6. Security Considerations
- 7. IANA Considerations
- 8. References
 - 8.1. Normative References
 - 8.2. Informative References

Acknowledgements

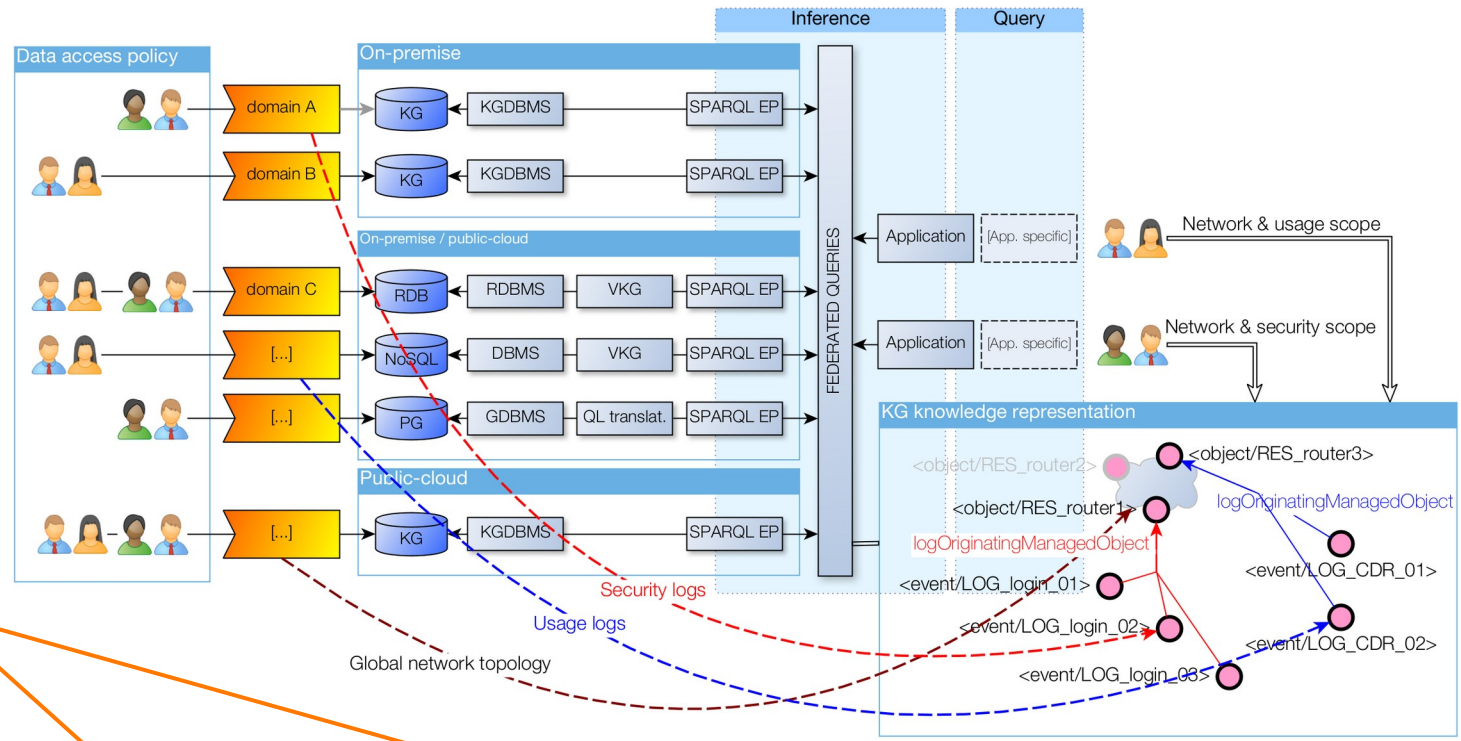


draft-tailhardat-nmop-incident-management-noria: overview -- §4.3.2

Table of Contents

- 1. Introduction
- 2. Conventions and Definitions
- 3. An ITSM-KG for Learning and Sharing Network Behavioral Models
 - 3.1. Principles
 - 3.2. Relation to the Digital Map
 - 3.2.1. Core Requirements
 - 3.2.2. Design Requirements
 - 3.2.3. Architectural Requirements
- 4. Strategies for the ITSM-KG Construction
 - 4.1. From YANG-based Configurations to Meta-Knowledge Graph
 - 4.2. Implementing Alignments of Model-Specificities to a Multi-Faceted Knowledge Graph
 - 4.2.1. The Network of Ontologies Approach
 - 4.2.2. Explicit Linking in the ONTO-META
 - 4.3. Extract-Transform-Load Pipelines for the ITSM-KG
 - 4.3.1. Handling Event Streams
 - 4.3.2. Federated Data Architecture
- 5. Experiments
 - 5.1. Experimental Plan
 - 5.2. Implementation Status
 - 5.2.1. NORIA
 - 5.2.2. YANG2OWL
- 6. Security Considerations
- 7. IANA Considerations
- 8. References
 - 8.1. Normative References
 - 8.2. Informative References

Acknowledgments



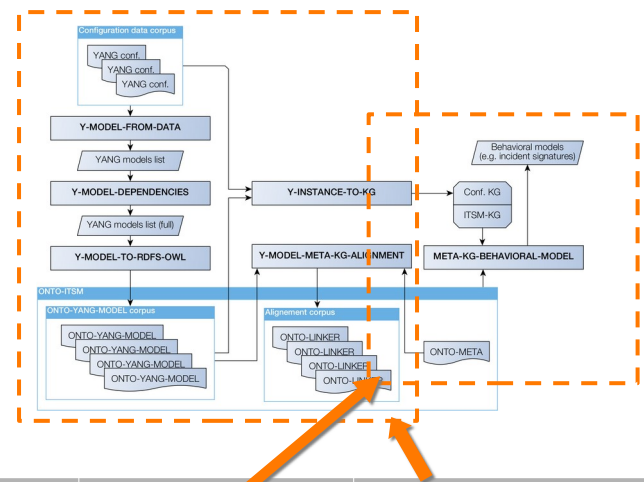
Federated querying over a multi-siloed graph
Providing **unified access** to data distributed across various technological platforms and stakeholders thanks to SPARQL Federated Queries and the use of a shared ONTO-ITSM across data management platforms.

draft-tailhardat-nmop-incident-management-noria: overview -- §5.2

Table of Contents

- 1. Introduction
- 2. Conventions and Definitions
- 3. An ITSM-KG for Learning and Sharing Network Behavioral Models
 - 3.1. Principles
 - 3.2. Relation to the Digital Map
 - 3.2.1. Core Requirements
 - 3.2.2. Design Requirements
 - 3.2.3. Architectural Requirements
- 4. Strategies for the ITSM-KG Construction
 - 4.1. From YANG-based Configurations to Meta-Knowledge Graph
 - 4.2. Implementing Alignments of Model-Specificities to a Multi-Faceted Knowledge Graph
 - 4.2.1. The Network of Ontologies Approach
 - 4.2.2. Explicit Linking in the ONTO-META
 - 4.3. Extract-Transform-Load Pipelines for the ITSM-KG
 - 4.3.1. Handling Event Streams
 - 4.3.2. Federated Data Architecture
- 5. Experiments
 - 5.1. Experimental Plan
 - 5.2. Implementation Status
 - 5.2.1. NORIA
 - 5.2.2. YANG2OWL
- 6. Security Considerations
- 7. IANA Considerations
- 8. References
 - 8.1. Normative References
 - 8.2. Informative References
- Acknowledgements

Implementation status ...
Two complementary experiments available as for now.

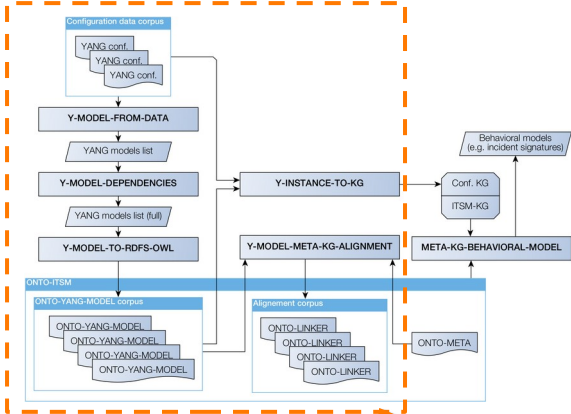


YANG-KG-SEMANTIC-GENERALIZATION use cases	NORIA draft v01	YANG2OWL draft v02
Y-MODEL-FROM-DATA	no	potential
Y-MODEL-DEPENDENCIES	no	yes
Y-MODEL-TO-RDFS-OWL	no	yes
Y-INSTANCE-TO-KG	potential	yes
Y-MODEL-META-KG-ALIGNMENT	potential	potential
META-KG-BEHAVIORAL-MODEL	yes	no

Focusing on the YANG2OWL approach ...

Table of Contents

- 1. Introduction
- 2. Conventions and Definitions
- 3. An ITSM-KG for Learning and Sharing Network Behavioral Models
 - 3.1. Principles
 - 3.2. Relation to the Digital Map
 - 3.2.1. Core Requirements
 - 3.2.2. Design Requirements
 - 3.2.3. Architectural Requirements
- 4. Strategies for the ITSM-KG Construction
 - 4.1. From YANG-based Configurations to Meta-Knowledge Graph
 - 4.2. Implementing Alignments of Model-Specificities to a Multi-Faceted Knowledge Graph
 - 4.2.1. The Network of Ontologies Approach
 - 4.2.2. Explicit Linking in the ONTO-META
 - 4.3. Extract-Transform-Load Pipelines for the ITSM-KG
 - 4.3.1. Handling Event Streams
 - 4.3.2. Federated Data Architecture
- 5. Experiments
 - 5.1. Experimental Plan
 - 5.2. Implementation Status
 - 5.2.1. NORIA
 - 5.2.2. YANG2OWL
- 6. Security Considerations
- 7. IANA Considerations
- 8. References
 - 8.1. Normative References
 - 8.2. Informative References
- Acknowledgments



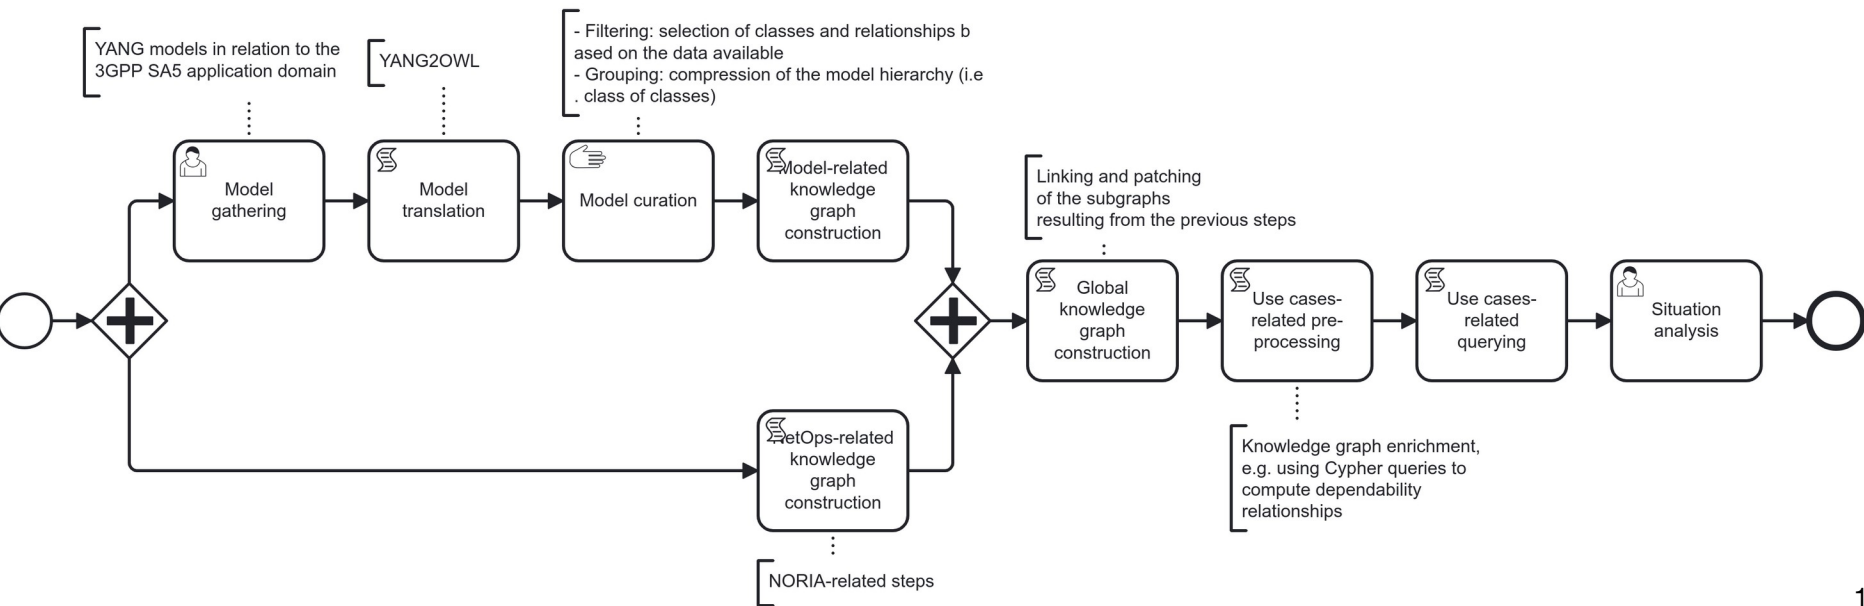
YANG-KG-SEMANTIC-GENERALIZATION use cases	NORIA draft v01	YANG2OWL draft v02
Y-MODEL-FROM-DATA	no	potential
Y-MODEL-DEPENDENCIES	no	yes
Y-MODEL-TO-RDFS-OWL	no	yes
Y-INSTANCE-TO-KG	potential	yes
Y-MODEL-META-KG-ALIGNMENT	potential	potential
META-KG-BEHAVIORAL-MODEL	yes	no

The YANG2OWL approach with an example of implementation -- §5.2.2.4

Data Virtualized 5G infrastructure (YANG based) + network ecosystem (other sources)

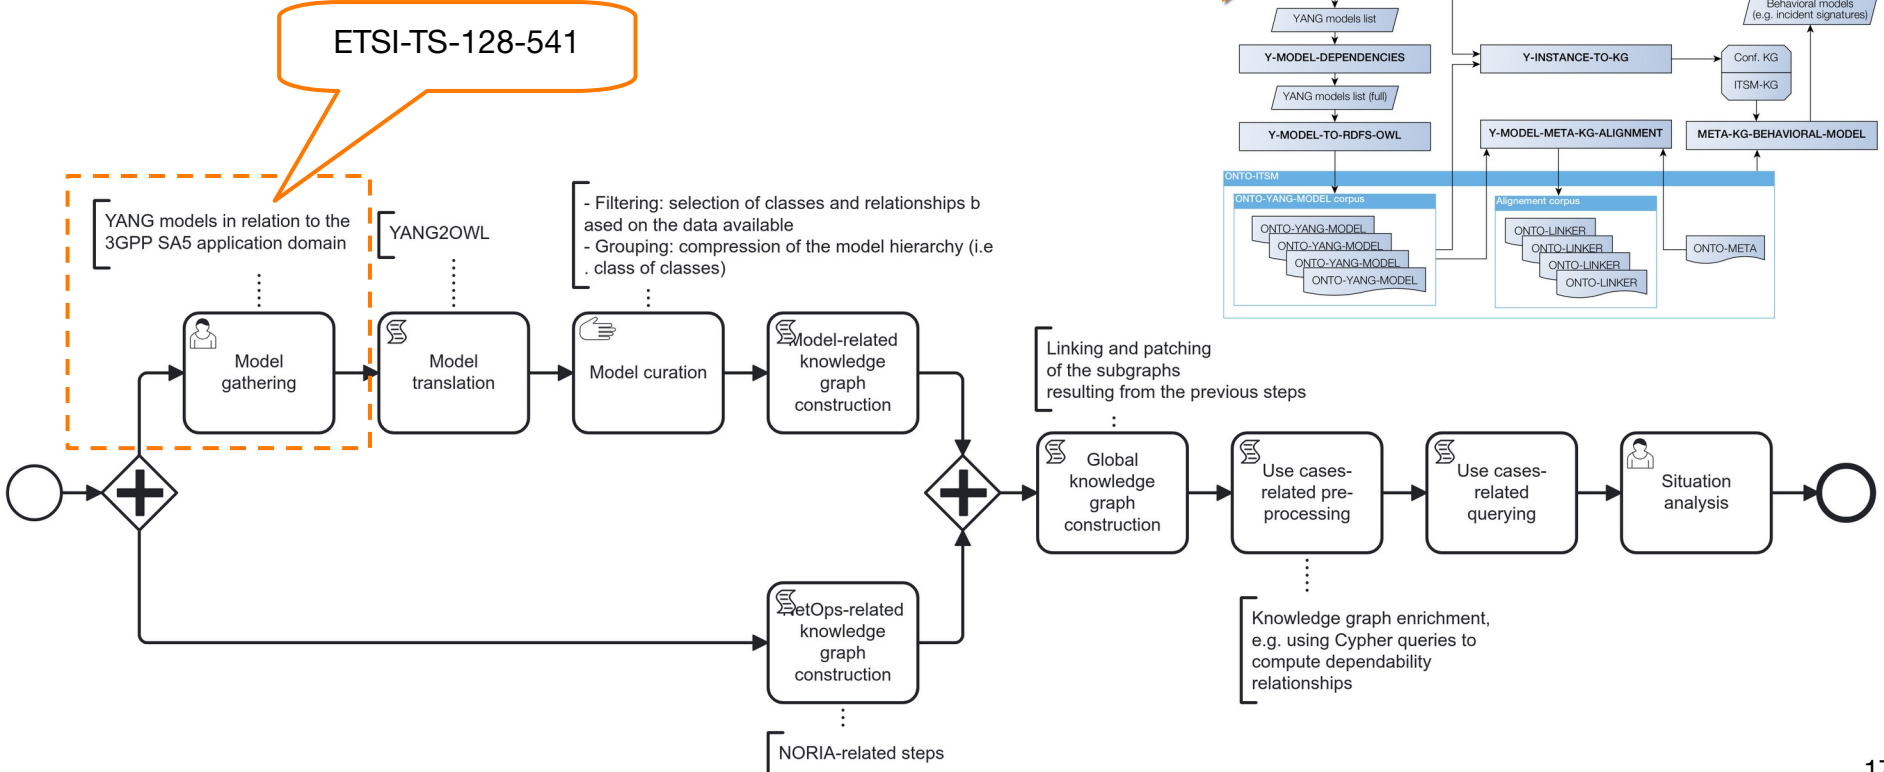
Operational context Network change management process -- impact analysis

Typical case For a scheduled operation on a leaf node (i.e. a network element in a 2-tier spine-leaf architecture), return all the servers connected to the leaf, all the Virtual Machines (VMs) hosted on these servers, all the Network Functions (NFs) deployed on these VMs, and ideally all the telecom services using these NFs.



The YANG2OWL approach with an example of implementation -- §5.2.2.4

ETSI-TS-128-541

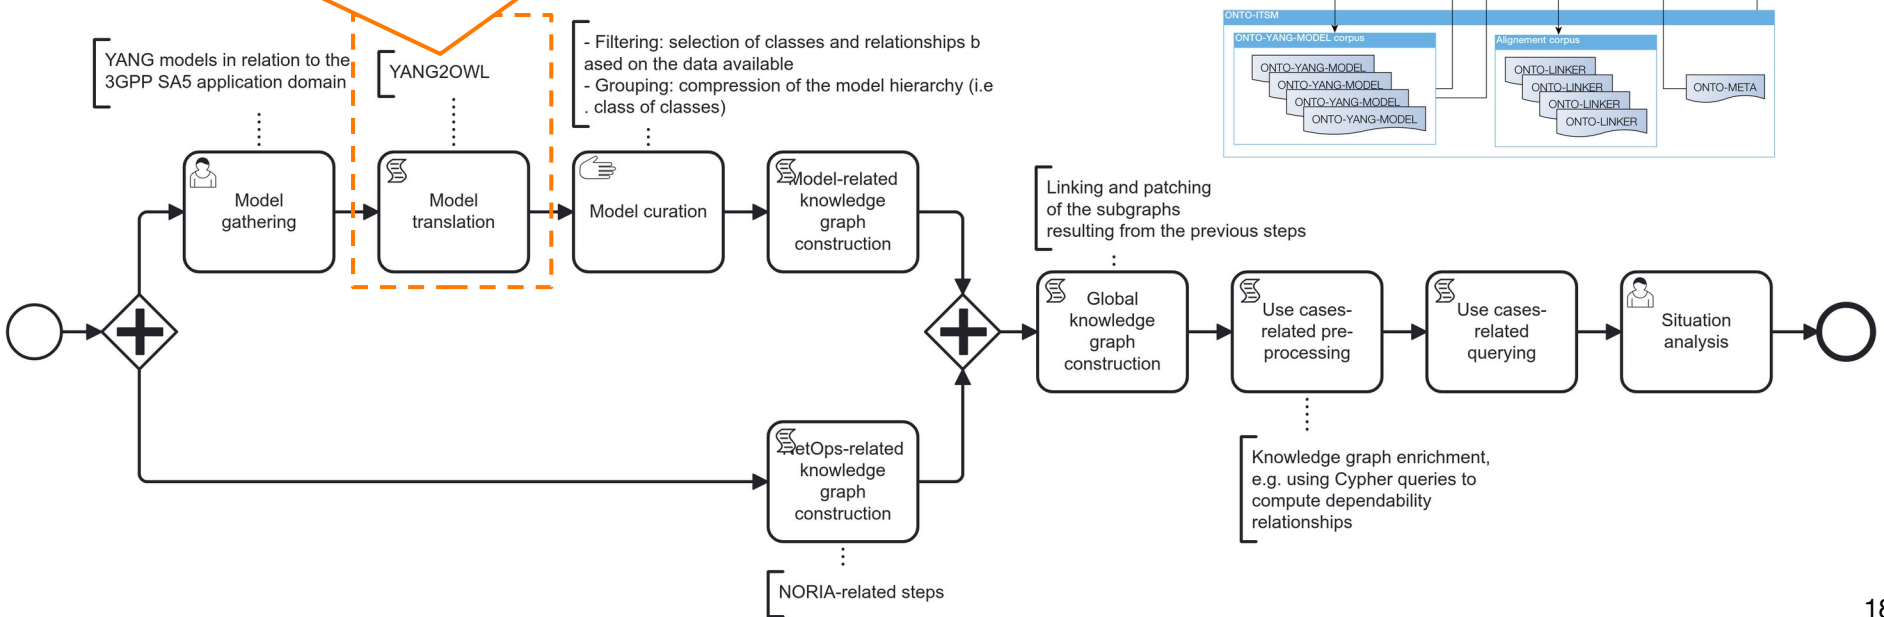


The YANG2OWL approach with an example of implementation -- §5.2.2.4

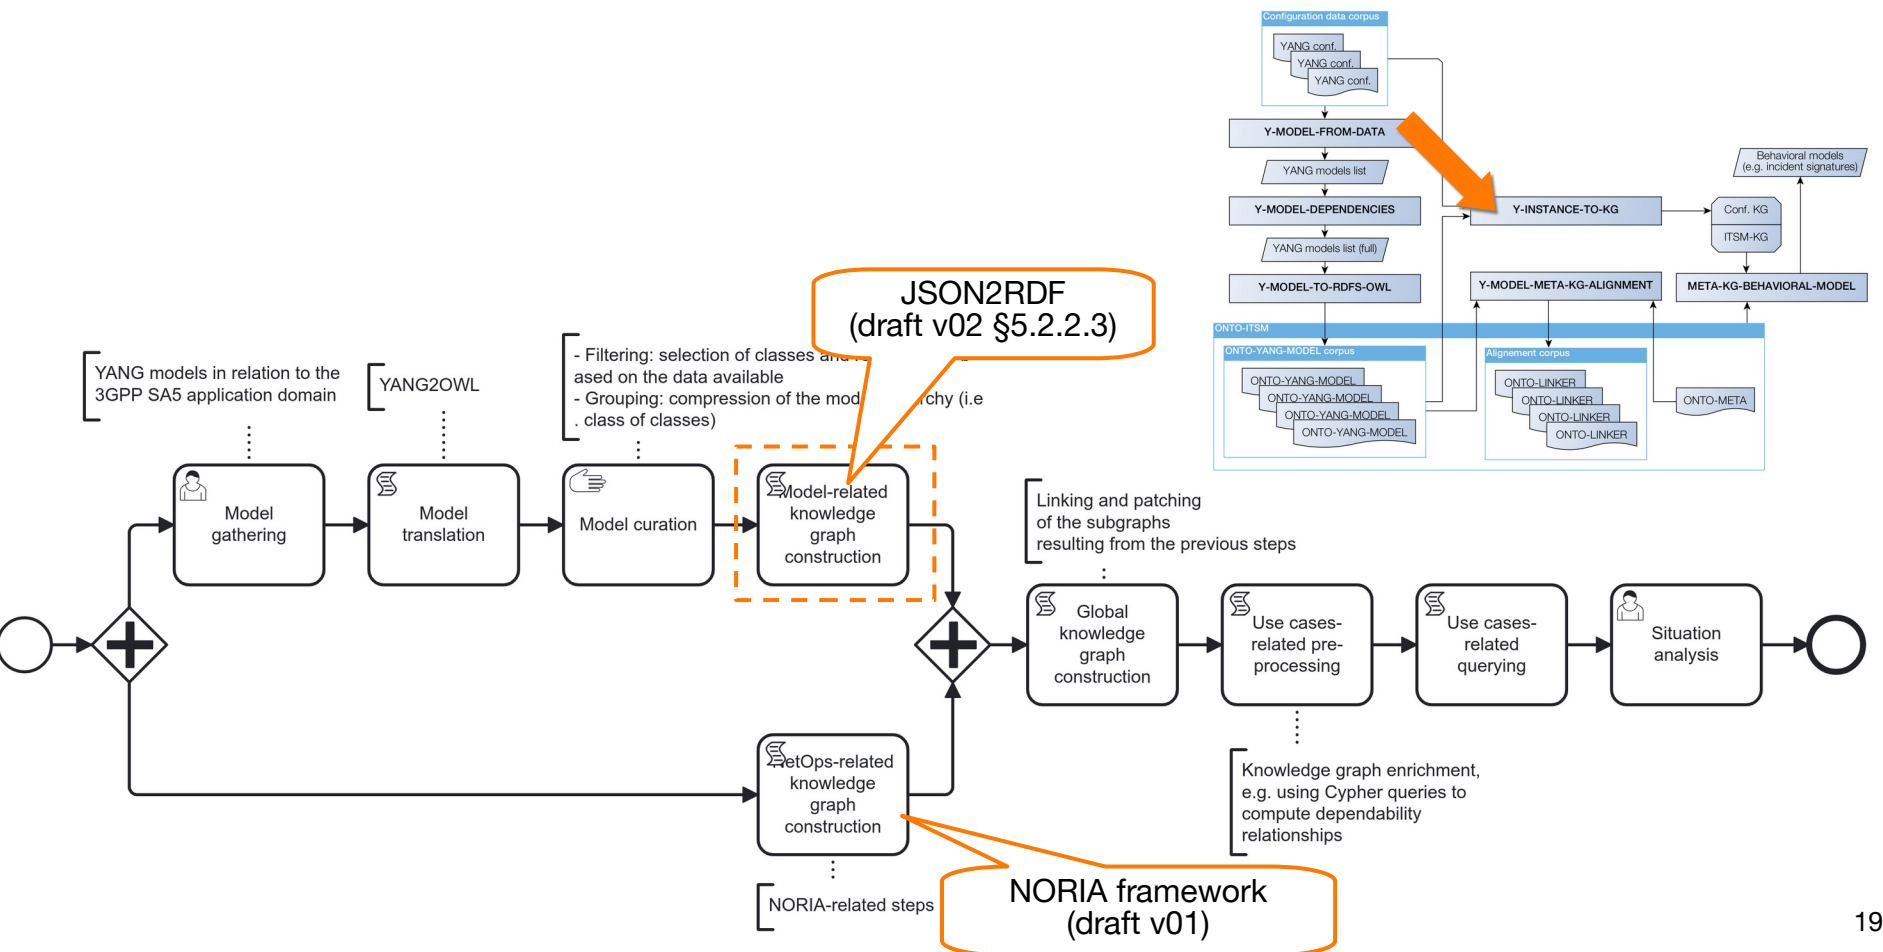
YANG2OWL converter (draft v02 §5.2.2.2)

Mapping rules between YANG constructs and OWL concepts :

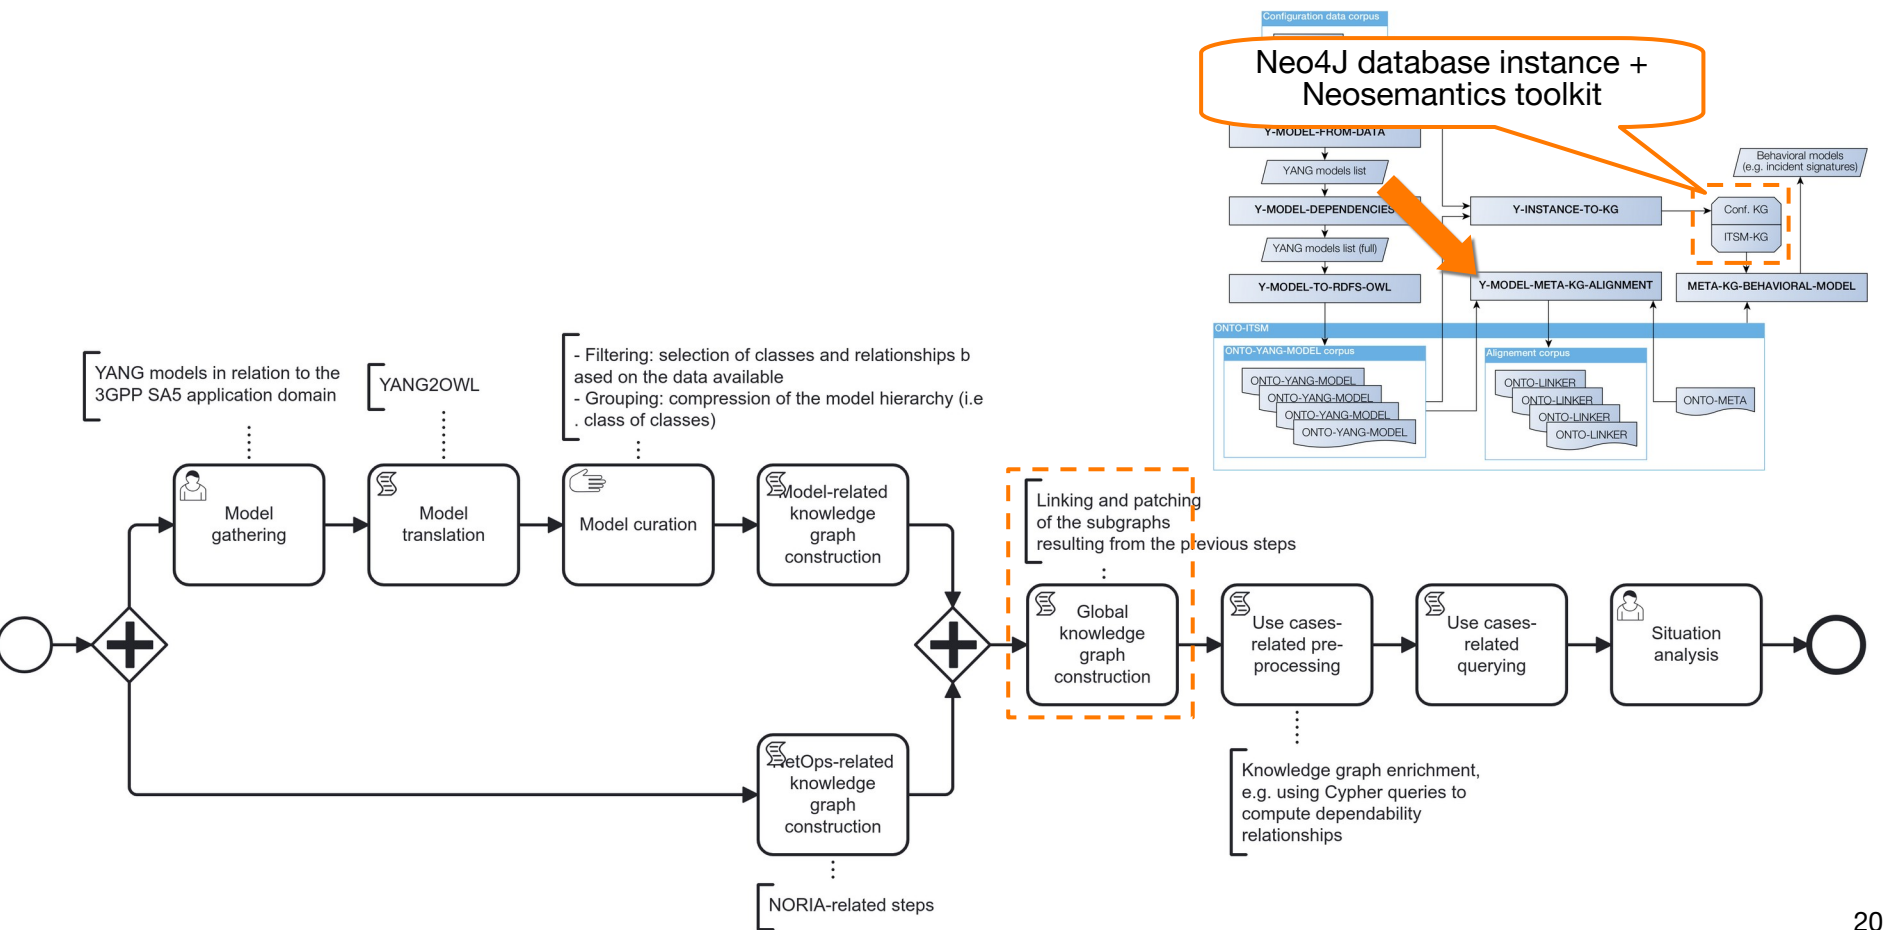
- **{container, list} → OWL classes.**
The name of the OWL class corresponds to the name of the container or list in the YANG model.
- **{leaf, leaf-list} → OWL data properties.**
The name of the OWL data property corresponds the name of the leaf or leaf-list in the YANG model.



The YANG2OWL approach with an example of implementation -- §5.2.2.4



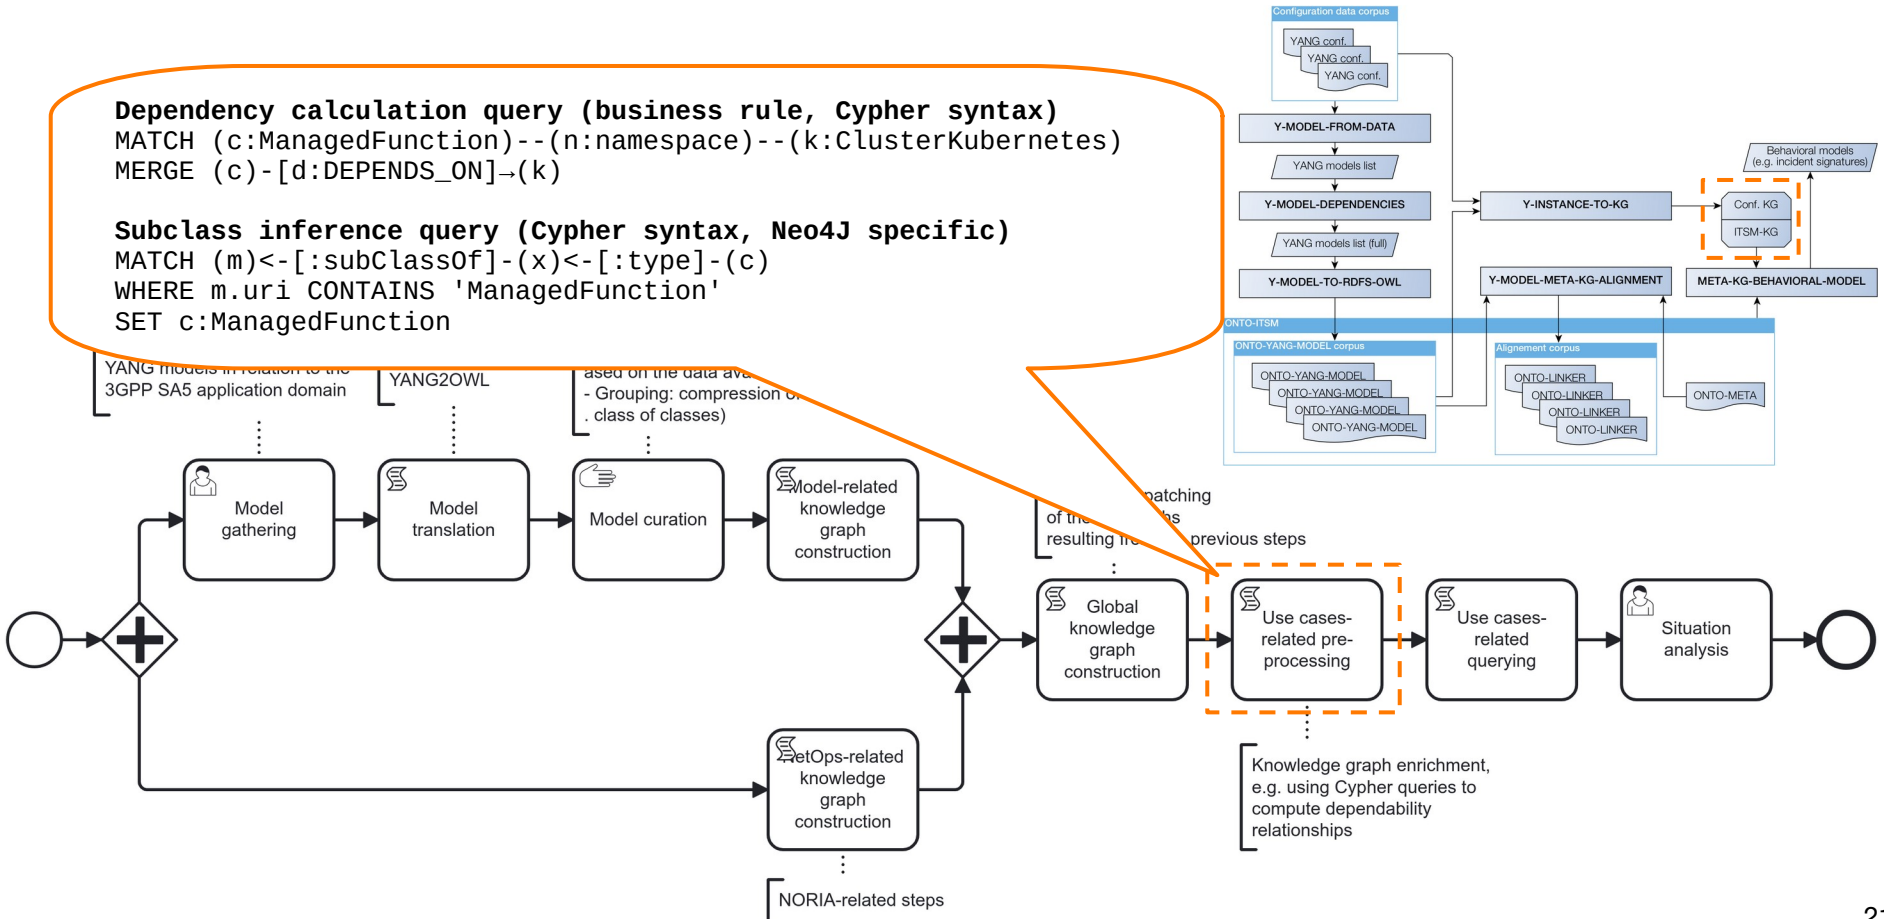
The YANG2OWL approach with an example of implementation -- §5.2.2.4



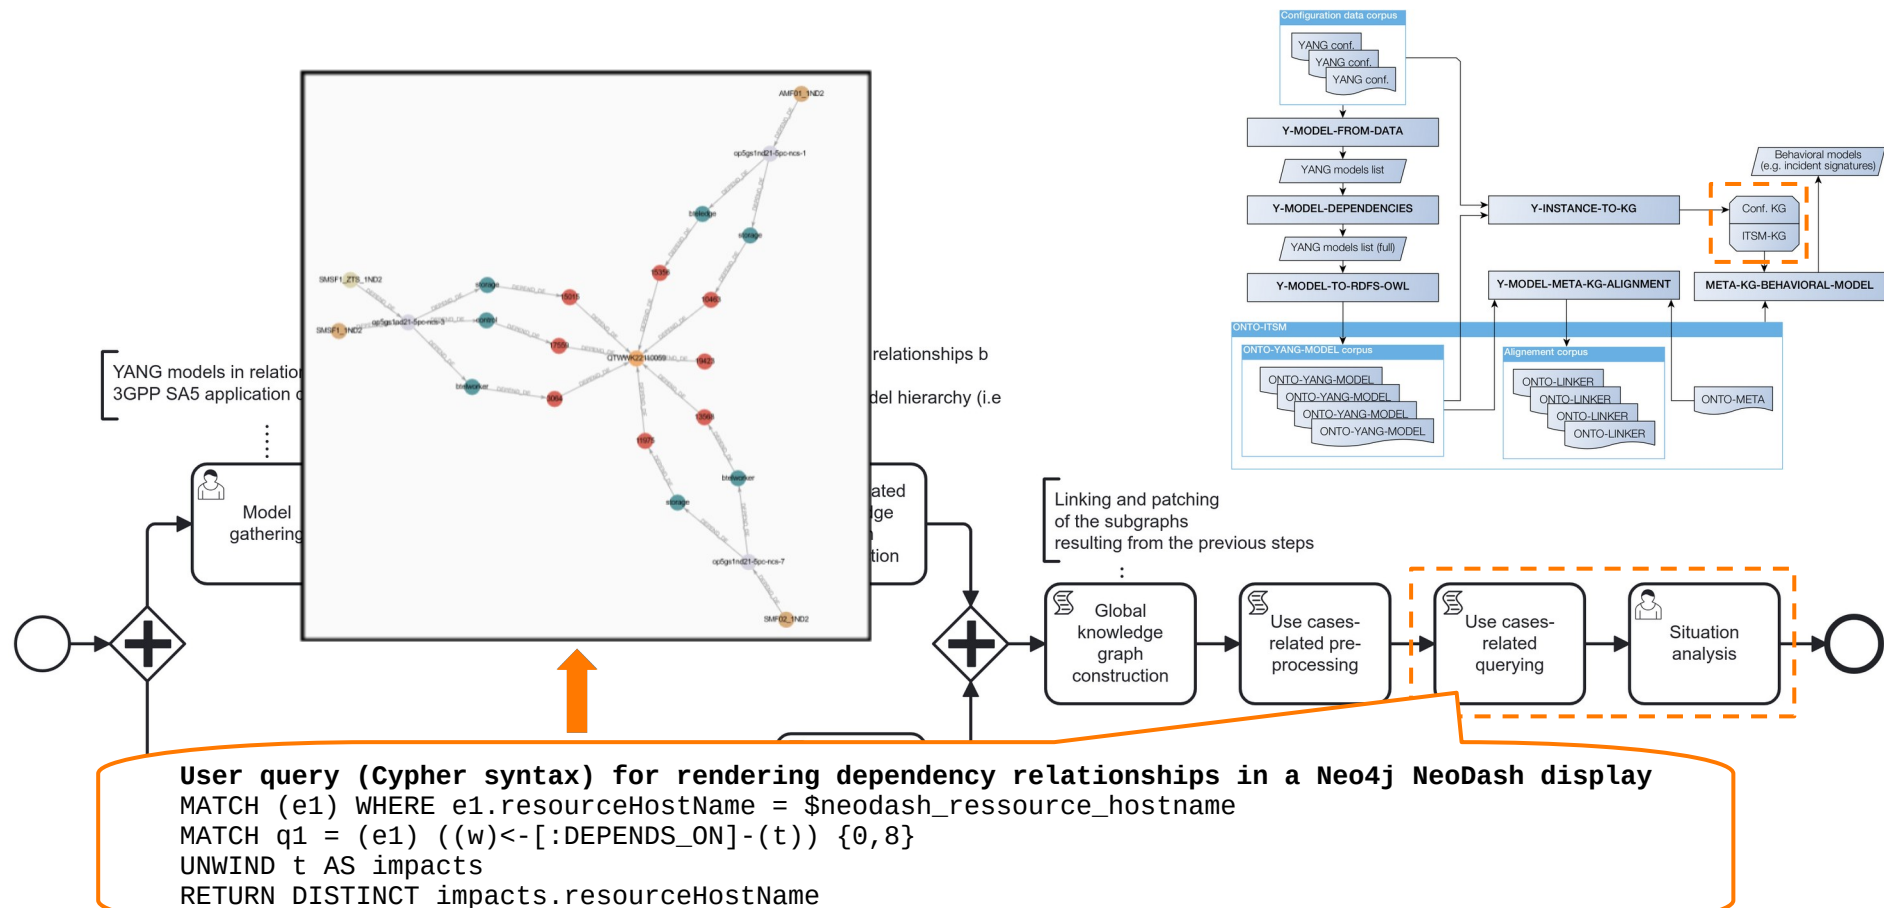
The YANG2OWL approach with an example of implementation -- §5.2.2.4

Dependency calculation query (business rule, Cypher syntax)
MATCH (c:ManagedFunction)--(n:namespace)--(k:ClusterKubernetes)
MERGE (c)-[d:DEPENDS_ON]-(k)

Subclass inference query (Cypher syntax, Neo4J specific)
MATCH (m)-[:subClassOf]-(x)-[:type]-(c)
WHERE m.uri CONTAINS 'ManagedFunction'
SET c:ManagedFunction



The YANG2OWL approach with an example of implementation -- §5.2.2.4



draft-tailhardat-nmop-incident-management-noria-02

<https://datatracker.ietf.org/doc/draft-tailhardat-nmop-incident-management-noria/>

Problem Building an ITSM Knowledge Graph that uses YANG-based configuration data while abstracting network details for learning and sharing behavioral models.

Approach Knowledge representation using SemWeb technologies, generalization of YANG models for configuration data, an extended Digital Map combining configuration with operational and OSS data, and a data processing pipeline for experimentation.

Next Call for experiments and contributions on the draft-tailhardat-nmop-incident-management-noria proposal.

YANG-KG-SEMANTIC-GENERALIZATION use cases	NORIA draft v01	YANG2OWL draft v02
Y-MODEL-FROM-DATA	no	potential
Y-MODEL-DEPENDENCIES	no	yes
Y-MODEL-TO-RDFS-OWL	no	yes
Y-INSTANCE-TO-KG	potential	yes
Y-MODEL-META-KG-ALIGNMENT	potential	potential
META-KG-BEHAVIORAL-MODEL	yes	no

draft-tailhardat-nmop-incident-management-noria-02

<https://datatracker.ietf.org/doc/draft-tailhardat-nmop-incident-management-noria/>

Problem Building an ITSM Knowledge Graph that uses YANG-based configuration data while abstracting network details for learning and sharing behavioral models.

Approach Knowledge representation using SemWeb technologies, generalization of YANG models for configuration data, an extended Digital Map combining configuration with operational and OSS data, and a data processing pipeline for experimentation.

Next Call for experiments and contributions on the draft-tailhardat-nmop-incident-management-noria proposal.

YANG-KG-SEMANTIC-GENERALIZATION use cases	NORIA draft v01	YANG2OWL draft v02
Y-MODEL-FROM-DATA	no	potential
Y-MODEL-DEPENDENCIES	no	yes
Y-MODEL-TO-RDFS-OWL	no	yes
Y-INSTANCE-TO-KG	potential	yes
Y-MODEL-META-KG-ALIGNMENT	potential	potential
META-KG-BEHAVIORAL-MODEL	yes	no

Implementation status in short ...

NORIA Means for building a **unified view of complex ICT systems** and learning/exploiting/sharing network behavioral models.

YANG2OWL Streamlines the **development of NDT architectures** based on knowledge graphs and simplifies ITSM-KG updates when YANG modules change. It notably automates the Ontology Implementation and Ontology Update activities of the LOT4KG methodology [LOT4KG-2024].


draft-tailhardat-nmop-incident-management-noria-02

<https://datatracker.ietf.org/doc/draft-tailhardat-nmop-incident-management-noria/>

Problem Building an ITSM Knowledge Graph that uses YANG-based configuration data while abstracting network details for learning and sharing behavioral models.

Approach Knowledge representation using SemWeb technologies, generalization of YANG models for configuration data, an extended Digital Map combining configuration with operational and OSS data, and a data processing pipeline for experimentation.

Next Call for experiments and contributions on the draft-tailhardat-nmop-incident-management-noria proposal.



YANG-KG-SEMANTIC-GENERALIZATION use cases

Y-MODEL-FROM-DATA

Y-MODEL-DEPENDENCIES

Y-MODEL-TO-RDFS-OWL

Y-INSTANCE-TO-KG

Y-MODEL-META-KG-ALIGNMENT

META-KG-BEHAVIORAL-MODELS

Ideas for new opportunities ...

- Learning and sharing anomaly models using the « **AnTagOnIst** » (Anomaly Tagging On Historical data) [GitHub] framework?
- Building the ITSM-KG with the « **Declarative Construction** of Knowledge Graphs from **NETCONF Data Sources** » (Dominguez, et al. - 2025) [SWJ] toolkit?
- Combine the « **YANG2RDF** » and « **YANG2OWL** » approaches?
- Reflect on how to **automate the Y-MODEL-META-KG-ALIGNMENT** use case.
- Check if there exists **universal YANG ⇔ RDFS/OWL translation** principles?
- Reflect on how to ensure **reliable retrieval of dependencies between YANG modules** for the Y-MODEL-DEPENDENCIES use case.